

Evaluating Diagnostic Models

January 5, 2024

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
import util
from public_tests import *
from test_utils import *
```

```
In [5]: train_results = pd.read_csv("data/train_preds.csv")
valid_results = pd.read_csv("data/valid_preds.csv")
```

```
# the labels in our dataset
class_labels = ['Cardiomegaly',
                'Emphysema',
                'Effusion',
                'Hernia',
                'Infiltration',
                'Mass',
                'Nodule',
                'Atelectasis',
                'Pneumothorax',
                'Pleural_Thickening',
                'Pneumonia',
                'Fibrosis',
                'Edema',
                'Consolidation']
```

```
# the labels for prediction values in our dataset
pred_labels = [l + "_pred" for l in class_labels]
```

```
In [6]: y = valid_results[class_labels].values
pred = valid_results[pred_labels].values
```

```
In [7]: # let's take a peek at our dataset
valid_results[np.concatenate([class_labels, pred_labels])].head()
```

```
Out[7]:   Cardiomegaly  Emphysema  Effusion  Hernia  Infiltration  Mass  Nodule  \
0              0              0          0        0              0      0        0
```

1	0	0	0	0	1	0	1
2	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0

	Atelectasis	Pneumothorax	Pleural_Thickening	...	Infiltration_pred	\
0	0	0	0	...	0.256020	
1	0	0	0	...	0.382199	
2	0	0	0	...	0.427727	
3	0	0	0	...	0.158596	
4	0	0	0	...	0.536762	


	Mass_pred	Nodule_pred	Atelectasis_pred	Pneumothorax_pred	\
0	0.266928	0.312440	0.460342	0.079453	
1	0.176825	0.465807	0.489424	0.084595	
2	0.115513	0.249030	0.035105	0.238761	
3	0.259460	0.334870	0.266489	0.073371	
4	0.198797	0.273110	0.186771	0.242122	

	Pleural_Thickening_pred	Pneumonia_pred	Fibrosis_pred	Edema_pred	\
0	0.271495	0.276861	0.398799	0.015867	
1	0.377318	0.363582	0.638024	0.025948	
2	0.167095	0.166389	0.262463	0.007758	
3	0.229834	0.191281	0.344348	0.008559	
4	0.309786	0.411771	0.244666	0.126930	

	Consolidation_pred
0	0.156320
1	0.144419
2	0.125790
3	0.119153
4	0.342409

[5 rows x 28 columns]

```
In [8]: plt.xticks(rotation=90)
plt.bar(x = class_labels, height= y.sum(axis=0));
```



```
In [9]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def true_positives(y, pred, th=0.5):
    """
    Count true positives.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        TP (int): true positives
    """
    TP = 0

    # get thresholded predictions
    thresholded_preds = pred >= th

    # compute TP
    TP = np.sum((y == 1) & (thresholded_preds == 1))
```

```

    return TP

def true_negatives(y, pred, th=0.5):
    """
    Count true negatives.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        TN (int): true negatives
    """
    TN = 0

    # get thresholded predictions
    thresholded_preds = pred >= th

    # compute TN
    TN = np.sum((y == 0) & (thresholded_preds == 0))

    return TN

def false_positives(y, pred, th=0.5):
    """
    Count false positives.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        FP (int): false positives
    """
    FP = 0

    # get thresholded predictions
    thresholded_preds = pred >= th

    # compute FP
    FP = np.sum((y == 0) & (thresholded_preds == 1))

    return FP

def false_negatives(y, pred, th=0.5):
    """
    Count false positives.

```

```

Args:
    y (np.array): ground truth, size (n_examples)
    pred (np.array): model output, size (n_examples)
    th (float): cutoff value for positive prediction from model
Returns:
    FN (int): false negatives
    """
    FN = 0

    # get thresholded predictions
    thresholded_preds = pred >= th

    # compute FN
    FN = np.sum((y == 1) & (thresholded_preds == 0))

    return FN

```

```

In [10]: ### do not modify this cell
         get_tp_tn_fp_fn_test(true_positives, true_negatives, false_positives, false_negatives)

```

	y_test	preds_test	category
0	1	0.8	TP
1	1	0.7	TP
2	0	0.4	TN
3	0	0.3	TN
4	0	0.2	TN
5	0	0.5	FP
6	0	0.6	FP
7	0	0.7	FP
8	0	0.8	FP
9	1	0.1	FN
10	1	0.2	FN
11	1	0.3	FN
12	1	0.4	FN
13	1	0.0	FN

Your functions calculated:

```

TP: 2
TN: 3
FP: 4
FN: 5

```

All tests passed. All tests passed. All tests passed. All tests passed.

```

In [ ]: util.get_performance_metrics(y, pred, class_labels)

```

```

In [11]: # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
         def get_accuracy(y, pred, th=0.5):

```

```

"""
Compute accuracy of predictions at threshold.

Args:
    y (np.array): ground truth, size (n_examples)
    pred (np.array): model output, size (n_examples)
    th (float): cutoff value for positive prediction from model
Returns:
    accuracy (float): accuracy of predictions at threshold
"""
accuracy = 0.0

### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

# get TP, FP, TN, FN using our previously defined functions
TP = true_positives(y, pred, th)
FP = false_positives(y, pred, th)
TN = true_negatives(y, pred, th)
FN = false_negatives(y, pred, th)

# Compute accuracy using TP, FP, TN, FN
accuracy = (TP + TN) / (TP + TN + FP + FN) if (TP + TN + FP + FN) != 0 else 0.0

### END CODE HERE ###

return accuracy

```

```

In [12]: ### do not modify this cell
         get_accuracy_test(get_accuracy)

```

Test Case:

```

Test Labels:      [1 0 0 1 1]
Test Predictions: [0.8 0.8 0.4 0.6 0.3]
Threshold:        0.5
Computed Accuracy: 0.6

```

All tests passed.

```

In [13]: util.get_performance_metrics(y, pred, class_labels, acc=get_accuracy)

```

```

Out[13]:

```

	TP	TN	FP	FN	Accuracy	Prevalence	Sensitivity \
Cardiomegaly	16	814	169	1	0.83	Not Defined	Not Defined
Emphysema	20	869	103	8	0.889	Not Defined	Not Defined
Effusion	99	690	196	15	0.789	Not Defined	Not Defined
Hernia	1	743	255	1	0.744	Not Defined	Not Defined
Infiltration	114	543	265	78	0.657	Not Defined	Not Defined
Mass	40	789	158	13	0.829	Not Defined	Not Defined

Nodule	28	731	220	21	0.759	Not Defined	Not Defined
Atelectasis	64	657	249	30	0.721	Not Defined	Not Defined
Pneumothorax	24	785	183	8	0.809	Not Defined	Not Defined
Pleural_Thickening	24	713	259	4	0.737	Not Defined	Not Defined
Pneumonia	14	661	320	5	0.675	Not Defined	Not Defined
Fibrosis	10	725	261	4	0.735	Not Defined	Not Defined
Edema	15	767	213	5	0.782	Not Defined	Not Defined
Consolidation	36	658	297	9	0.694	Not Defined	Not Defined

	Specificity	PPV	NPV	AUC \
Cardiomegaly	Not Defined	Not Defined	Not Defined	Not Defined
Emphysema	Not Defined	Not Defined	Not Defined	Not Defined
Effusion	Not Defined	Not Defined	Not Defined	Not Defined
Hernia	Not Defined	Not Defined	Not Defined	Not Defined
Infiltration	Not Defined	Not Defined	Not Defined	Not Defined
Mass	Not Defined	Not Defined	Not Defined	Not Defined
Nodule	Not Defined	Not Defined	Not Defined	Not Defined
Atelectasis	Not Defined	Not Defined	Not Defined	Not Defined
Pneumothorax	Not Defined	Not Defined	Not Defined	Not Defined
Pleural_Thickening	Not Defined	Not Defined	Not Defined	Not Defined
Pneumonia	Not Defined	Not Defined	Not Defined	Not Defined
Fibrosis	Not Defined	Not Defined	Not Defined	Not Defined
Edema	Not Defined	Not Defined	Not Defined	Not Defined
Consolidation	Not Defined	Not Defined	Not Defined	Not Defined

F1 Threshold

Cardiomegaly	Not Defined	0.5
Emphysema	Not Defined	0.5
Effusion	Not Defined	0.5
Hernia	Not Defined	0.5
Infiltration	Not Defined	0.5
Mass	Not Defined	0.5
Nodule	Not Defined	0.5
Atelectasis	Not Defined	0.5
Pneumothorax	Not Defined	0.5
Pleural_Thickening	Not Defined	0.5
Pneumonia	Not Defined	0.5
Fibrosis	Not Defined	0.5
Edema	Not Defined	0.5
Consolidation	Not Defined	0.5

```
In [14]: get_accuracy(valid_results["Emphysema"].values, np.zeros(len(valid_results)))
```

```
Out[14]: 0.972
```

```
In [15]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def get_prevalence(y):
```

```

"""
Compute prevalence.

Args:
    y (np.array): ground truth, size (n_examples)
Returns:
    prevalence (float): prevalence of positive cases
"""
prevalence = 0.0

### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

prevalence = np.mean(y)

### END CODE HERE ###

return prevalence

```

```

In [16]: ### do npt modify this cell
         get_prevalence_test(get_prevalence)

```

Test Case:

```

Test Labels:          [1 0 0 1 1 0 0 0 0 1]
Computed Prevalence:  0.4

```

All tests passed.

```

In [17]: util.get_performance_metrics(y, pred, class_labels, acc=get_accuracy, prevalence=get_prevalence)

```

```

Out[17]:

```

	TP	TN	FP	FN	Accuracy	Prevalence	Sensitivity	\
Cardiomegaly	16	814	169	1	0.83	0.017	Not Defined	
Emphysema	20	869	103	8	0.889	0.028	Not Defined	
Effusion	99	690	196	15	0.789	0.114	Not Defined	
Hernia	1	743	255	1	0.744	0.002	Not Defined	
Infiltration	114	543	265	78	0.657	0.192	Not Defined	
Mass	40	789	158	13	0.829	0.053	Not Defined	
Nodule	28	731	220	21	0.759	0.049	Not Defined	
Atelectasis	64	657	249	30	0.721	0.094	Not Defined	
Pneumothorax	24	785	183	8	0.809	0.032	Not Defined	
Pleural_Thickening	24	713	259	4	0.737	0.028	Not Defined	
Pneumonia	14	661	320	5	0.675	0.019	Not Defined	
Fibrosis	10	725	261	4	0.735	0.014	Not Defined	
Edema	15	767	213	5	0.782	0.02	Not Defined	
Consolidation	36	658	297	9	0.694	0.045	Not Defined	
	Specificity				PPV	NPV	AUC	\

Cardiomegaly	Not Defined	Not Defined	Not Defined	Not Defined
Emphysema	Not Defined	Not Defined	Not Defined	Not Defined
Effusion	Not Defined	Not Defined	Not Defined	Not Defined
Hernia	Not Defined	Not Defined	Not Defined	Not Defined
Infiltration	Not Defined	Not Defined	Not Defined	Not Defined
Mass	Not Defined	Not Defined	Not Defined	Not Defined
Nodule	Not Defined	Not Defined	Not Defined	Not Defined
Atelectasis	Not Defined	Not Defined	Not Defined	Not Defined
Pneumothorax	Not Defined	Not Defined	Not Defined	Not Defined
Pleural_Thickening	Not Defined	Not Defined	Not Defined	Not Defined
Pneumonia	Not Defined	Not Defined	Not Defined	Not Defined
Fibrosis	Not Defined	Not Defined	Not Defined	Not Defined
Edema	Not Defined	Not Defined	Not Defined	Not Defined
Consolidation	Not Defined	Not Defined	Not Defined	Not Defined

F1 Threshold

Cardiomegaly	Not Defined	0.5
Emphysema	Not Defined	0.5
Effusion	Not Defined	0.5
Hernia	Not Defined	0.5
Infiltration	Not Defined	0.5
Mass	Not Defined	0.5
Nodule	Not Defined	0.5
Atelectasis	Not Defined	0.5
Pneumothorax	Not Defined	0.5
Pleural_Thickening	Not Defined	0.5
Pneumonia	Not Defined	0.5
Fibrosis	Not Defined	0.5
Edema	Not Defined	0.5
Consolidation	Not Defined	0.5

```
In [18]: # UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
```

```
def get_sensitivity(y, pred, th=0.5):
    """
    Compute sensitivity of predictions at threshold.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        sensitivity (float): probability that our test outputs positive given that th
    """
    sensitivity = 0.0

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
```

```

    # get TP and FN using our previously defined functions
    TP = true_positives(y, pred, th)
    FN = false_negatives(y, pred, th)

    # use TP and FN to compute sensitivity
    sensitivity = TP / (TP + FN) if (TP + FN) != 0 else 0.0

    ### END CODE HERE ###

    return sensitivity

def get_specificity(y, pred, th=0.5):
    """
    Compute specificity of predictions at threshold.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        specificity (float): probability that the test outputs negative given that th
    """
    specificity = 0.0

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

    # get TN and FP using our previously defined functions
    TN = true_negatives(y, pred, th)
    FP = false_positives(y, pred, th)

    # use TN and FP to compute specificity
    specificity = TN / (TN + FP) if (TN + FP) != 0 else 0.0

    ### END CODE HERE ###

    return specificity

```

```

In [19]: ### do not modify this cell
         get_sensitivity_specificity_test(get_sensitivity, get_specificity)

```

Test Case:

```

Test Labels:           [1 0 0 1 1]
Test Predictions:      [1 0 0 1 1]
Threshold:             0.5
Computed Sensitivity:   0.6666666666666666
Computed Specificity:   0.5

```

All tests passed. All tests passed.

```
In [20]: util.get_performance_metrics(y, pred, class_labels, acc=get_accuracy, prevalence=get_prevalence,
sens=get_sensitivity, spec=get_specificity)
```

```
Out[20]:
```

	TP	TN	FP	FN	Accuracy	Prevalence	Sensitivity	\
--	----	----	----	----	----------	------------	-------------	---

Cardiomegaly	16	814	169	1	0.83	0.017	0.941	
Emphysema	20	869	103	8	0.889	0.028	0.714	
Effusion	99	690	196	15	0.789	0.114	0.868	
Hernia	1	743	255	1	0.744	0.002	0.5	
Infiltration	114	543	265	78	0.657	0.192	0.594	
Mass	40	789	158	13	0.829	0.053	0.755	
Nodule	28	731	220	21	0.759	0.049	0.571	
Atelectasis	64	657	249	30	0.721	0.094	0.681	
Pneumothorax	24	785	183	8	0.809	0.032	0.75	
Pleural_Thickening	24	713	259	4	0.737	0.028	0.857	
Pneumonia	14	661	320	5	0.675	0.019	0.737	
Fibrosis	10	725	261	4	0.735	0.014	0.714	
Edema	15	767	213	5	0.782	0.02	0.75	
Consolidation	36	658	297	9	0.694	0.045	0.8	

	Specificity	PPV	NPV	AUC	\
--	-------------	-----	-----	-----	---

Cardiomegaly	0.828	Not Defined	Not Defined	Not Defined	
Emphysema	0.894	Not Defined	Not Defined	Not Defined	
Effusion	0.779	Not Defined	Not Defined	Not Defined	
Hernia	0.744	Not Defined	Not Defined	Not Defined	
Infiltration	0.672	Not Defined	Not Defined	Not Defined	
Mass	0.833	Not Defined	Not Defined	Not Defined	
Nodule	0.769	Not Defined	Not Defined	Not Defined	
Atelectasis	0.725	Not Defined	Not Defined	Not Defined	
Pneumothorax	0.811	Not Defined	Not Defined	Not Defined	
Pleural_Thickening	0.734	Not Defined	Not Defined	Not Defined	
Pneumonia	0.674	Not Defined	Not Defined	Not Defined	
Fibrosis	0.735	Not Defined	Not Defined	Not Defined	
Edema	0.783	Not Defined	Not Defined	Not Defined	
Consolidation	0.689	Not Defined	Not Defined	Not Defined	

F1 Threshold

Cardiomegaly	Not Defined	0.5
Emphysema	Not Defined	0.5
Effusion	Not Defined	0.5
Hernia	Not Defined	0.5
Infiltration	Not Defined	0.5
Mass	Not Defined	0.5
Nodule	Not Defined	0.5
Atelectasis	Not Defined	0.5
Pneumothorax	Not Defined	0.5

Pleural_Thickening	Not Defined	0.5
Pneumonia	Not Defined	0.5
Fibrosis	Not Defined	0.5
Edema	Not Defined	0.5
Consolidation	Not Defined	0.5

```
In [21]: # UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def get_ppv(y, pred, th=0.5):
    """
    Compute PPV of predictions at threshold.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        PPV (float): positive predictive value of predictions at threshold
    """
    PPV = 0.0

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

    # get TP and FP using our previously defined functions
    TP = true_positives(y, pred, th)
    FP = false_positives(y, pred, th)

    # use TP and FP to compute PPV
    PPV = TP / (TP + FP) if (TP + FP) != 0 else 0.0

    ### END CODE HERE ###

    return PPV

def get_npv(y, pred, th=0.5):
    """
    Compute NPV of predictions at threshold.

    Args:
        y (np.array): ground truth, size (n_examples)
        pred (np.array): model output, size (n_examples)
        th (float): cutoff value for positive prediction from model
    Returns:
        NPV (float): negative predictive value of predictions at threshold
    """
    NPV = 0.0

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###
```

```

# get TN and FN using our previously defined functions
TN = true_negatives(y, pred, th)
FN = false_negatives(y, pred, th)

# use TN and FN to compute NPV
NPV = TN / (TN + FN) if (TN + FN) != 0 else 0.0

### END CODE HERE ###

return NPV

```

```

In [22]: ### do not modify this cell
         get_ppv_npv_test(get_ppv, get_npv)

```

Test Case:

```

Test Labels:      [1 0 0 1 1]
Test Predictions: [1 0 0 1 1]
Threshold:        0.5
Computed PPV:     0.6666666666666666
Computed NPV:     0.5

```

All tests passed. All tests passed.

```

In [23]: util.get_performance_metrics(y, pred, class_labels, acc=get_accuracy, prevalence=get_prevalence,
        sens=get_sensitivity, spec=get_specificity, ppv=get_ppv, npv=get_npv)

```

```

Out[23]:

```

	TP	TN	FP	FN	Accuracy	Prevalence	Sensitivity	\
Cardiomegaly	16	814	169	1	0.83	0.017	0.941	
Emphysema	20	869	103	8	0.889	0.028	0.714	
Effusion	99	690	196	15	0.789	0.114	0.868	
Hernia	1	743	255	1	0.744	0.002	0.5	
Infiltration	114	543	265	78	0.657	0.192	0.594	
Mass	40	789	158	13	0.829	0.053	0.755	
Nodule	28	731	220	21	0.759	0.049	0.571	
Atelectasis	64	657	249	30	0.721	0.094	0.681	
Pneumothorax	24	785	183	8	0.809	0.032	0.75	
Pleural_Thickening	24	713	259	4	0.737	0.028	0.857	
Pneumonia	14	661	320	5	0.675	0.019	0.737	
Fibrosis	10	725	261	4	0.735	0.014	0.714	
Edema	15	767	213	5	0.782	0.02	0.75	
Consolidation	36	658	297	9	0.694	0.045	0.8	


	Specificity	PPV	NPV	AUC	F1	\
Cardiomegaly	0.828	0.086	0.999	Not Defined	Not Defined	
Emphysema	0.894	0.163	0.991	Not Defined	Not Defined	
Effusion	0.779	0.336	0.979	Not Defined	Not Defined	

Hernia	0.744	0.004	0.999	Not Defined	Not Defined
Infiltration	0.672	0.301	0.874	Not Defined	Not Defined
Mass	0.833	0.202	0.984	Not Defined	Not Defined
Nodule	0.769	0.113	0.972	Not Defined	Not Defined
Atelectasis	0.725	0.204	0.956	Not Defined	Not Defined
Pneumothorax	0.811	0.116	0.99	Not Defined	Not Defined
Pleural_Thickening	0.734	0.085	0.994	Not Defined	Not Defined
Pneumonia	0.674	0.042	0.992	Not Defined	Not Defined
Fibrosis	0.735	0.037	0.995	Not Defined	Not Defined
Edema	0.783	0.066	0.994	Not Defined	Not Defined
Consolidation	0.689	0.108	0.987	Not Defined	Not Defined

Threshold

Cardiomegaly	0.5
Emphysema	0.5
Effusion	0.5
Hernia	0.5
Infiltration	0.5
Mass	0.5
Nodule	0.5
Atelectasis	0.5
Pneumothorax	0.5
Pleural_Thickening	0.5
Pneumonia	0.5
Fibrosis	0.5
Edema	0.5
Consolidation	0.5

In [24]: util.get_curve(y, pred, class_labels)



```
In [25]: from sklearn.metrics import roc_auc_score
util.get_performance_metrics(y, pred, class_labels, acc=get_accuracy, prevalence=get_prevalence,
                             sens=get_sensitivity, spec=get_specificity, ppv=get_ppv, npv=get_npv)
```

Out [25]:

	TP	TN	FP	FN	Accuracy	Prevalence	Sensitivity	\
--	----	----	----	----	----------	------------	-------------	---

Cardiomegaly	16	814	169	1	0.83	0.017	0.941	
Emphysema	20	869	103	8	0.889	0.028	0.714	
Effusion	99	690	196	15	0.789	0.114	0.868	
Hernia	1	743	255	1	0.744	0.002	0.5	
Infiltration	114	543	265	78	0.657	0.192	0.594	
Mass	40	789	158	13	0.829	0.053	0.755	
Nodule	28	731	220	21	0.759	0.049	0.571	
Atelectasis	64	657	249	30	0.721	0.094	0.681	
Pneumothorax	24	785	183	8	0.809	0.032	0.75	
Pleural_Thickening	24	713	259	4	0.737	0.028	0.857	
Pneumonia	14	661	320	5	0.675	0.019	0.737	
Fibrosis	10	725	261	4	0.735	0.014	0.714	
Edema	15	767	213	5	0.782	0.02	0.75	
Consolidation	36	658	297	9	0.694	0.045	0.8	

	Specificity	PPV	NPV	AUC	F1	Threshold
Cardiomegaly	0.828	0.086	0.999	0.933	Not Defined	0.5

Emphysema	0.894	0.163	0.991	0.935	Not Defined	0.5
Effusion	0.779	0.336	0.979	0.891	Not Defined	0.5
Hernia	0.744	0.004	0.999	0.644	Not Defined	0.5
Infiltration	0.672	0.301	0.874	0.696	Not Defined	0.5
Mass	0.833	0.202	0.984	0.888	Not Defined	0.5
Nodule	0.769	0.113	0.972	0.745	Not Defined	0.5
Atelectasis	0.725	0.204	0.956	0.781	Not Defined	0.5
Pneumothorax	0.811	0.116	0.99	0.826	Not Defined	0.5
Pleural_Thickening	0.734	0.085	0.994	0.868	Not Defined	0.5
Pneumonia	0.674	0.042	0.992	0.762	Not Defined	0.5
Fibrosis	0.735	0.037	0.995	0.801	Not Defined	0.5
Edema	0.783	0.066	0.994	0.856	Not Defined	0.5
Consolidation	0.689	0.108	0.987	0.799	Not Defined	0.5

```
In [26]: def bootstrap_auc(y, pred, classes, bootstraps = 100, fold_size = 1000):
    statistics = np.zeros((len(classes), bootstraps))

    for c in range(len(classes)):
        df = pd.DataFrame(columns=['y', 'pred'])
        df.loc[:, 'y'] = y[:, c]
        df.loc[:, 'pred'] = pred[:, c]
        # get positive examples for stratified sampling
        df_pos = df[df.y == 1]
        df_neg = df[df.y == 0]
        prevalence = len(df_pos) / len(df)
        for i in range(bootstraps):
            # stratified sampling of positive and negative examples
            pos_sample = df_pos.sample(n = int(fold_size * prevalence), replace=True)
            neg_sample = df_neg.sample(n = int(fold_size * (1-prevalence)), replace=True)

            y_sample = np.concatenate([pos_sample.y.values, neg_sample.y.values])
            pred_sample = np.concatenate([pos_sample.pred.values, neg_sample.pred.values])
            score = roc_auc_score(y_sample, pred_sample)
            statistics[c][i] = score

    return statistics

statistics = bootstrap_auc(y, pred, class_labels)
```

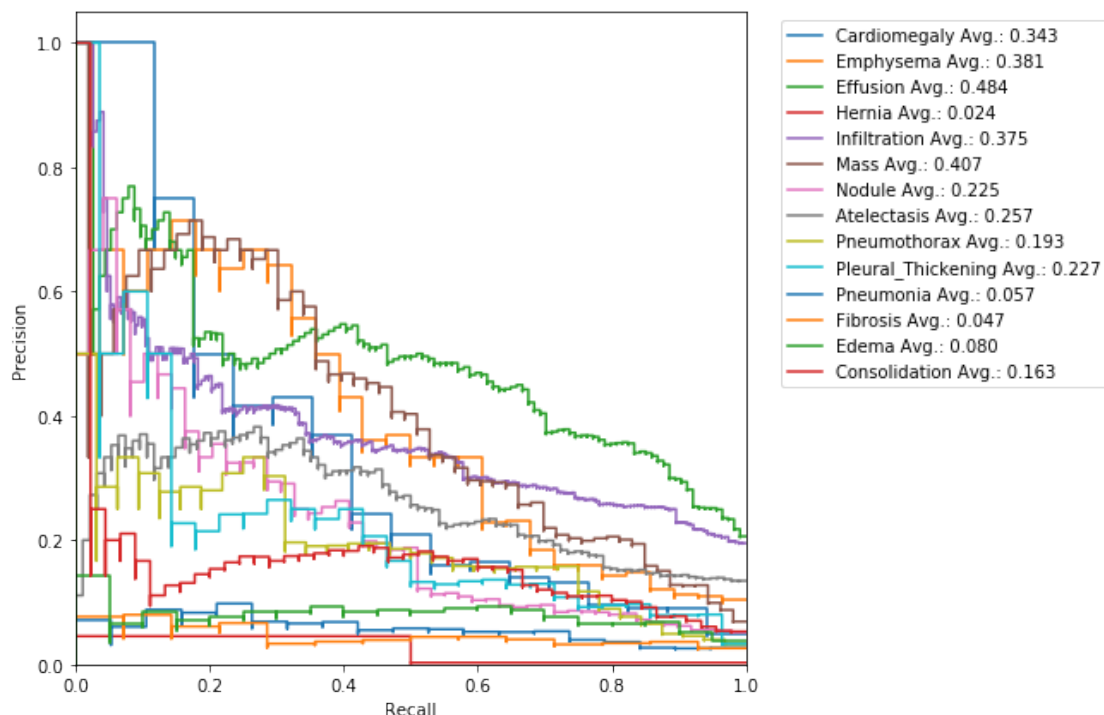
In [27]: util.print_confidence_intervals(class_labels, statistics)

```
Out[27]:
```

	Mean AUC (CI 5%-95%)
Cardiomegaly	0.93 (0.91-0.97)
Emphysema	0.94 (0.90-0.96)
Effusion	0.89 (0.87-0.91)
Hernia	0.64 (0.29-0.98)
Infiltration	0.70 (0.66-0.73)
Mass	0.89 (0.85-0.92)
Nodule	0.74 (0.69-0.81)

Atelectasis	0.78 (0.75-0.82)
Pneumothorax	0.83 (0.75-0.91)
Pleural_Thickening	0.86 (0.82-0.92)
Pneumonia	0.76 (0.67-0.83)
Fibrosis	0.80 (0.74-0.86)
Edema	0.86 (0.80-0.90)
Consolidation	0.80 (0.75-0.86)

In [28]: `util.get_curve(y, pred, class_labels, curve='prc')`



In [29]: `from sklearn.metrics import f1_score`
`util.get_performance_metrics(y, pred, class_labels, acc=get_accuracy, prevalence=get_prevalence,`
`sens=get_sensitivity, spec=get_specificity, ppv=get_ppv, npv=get_npv)`

Out [29]:

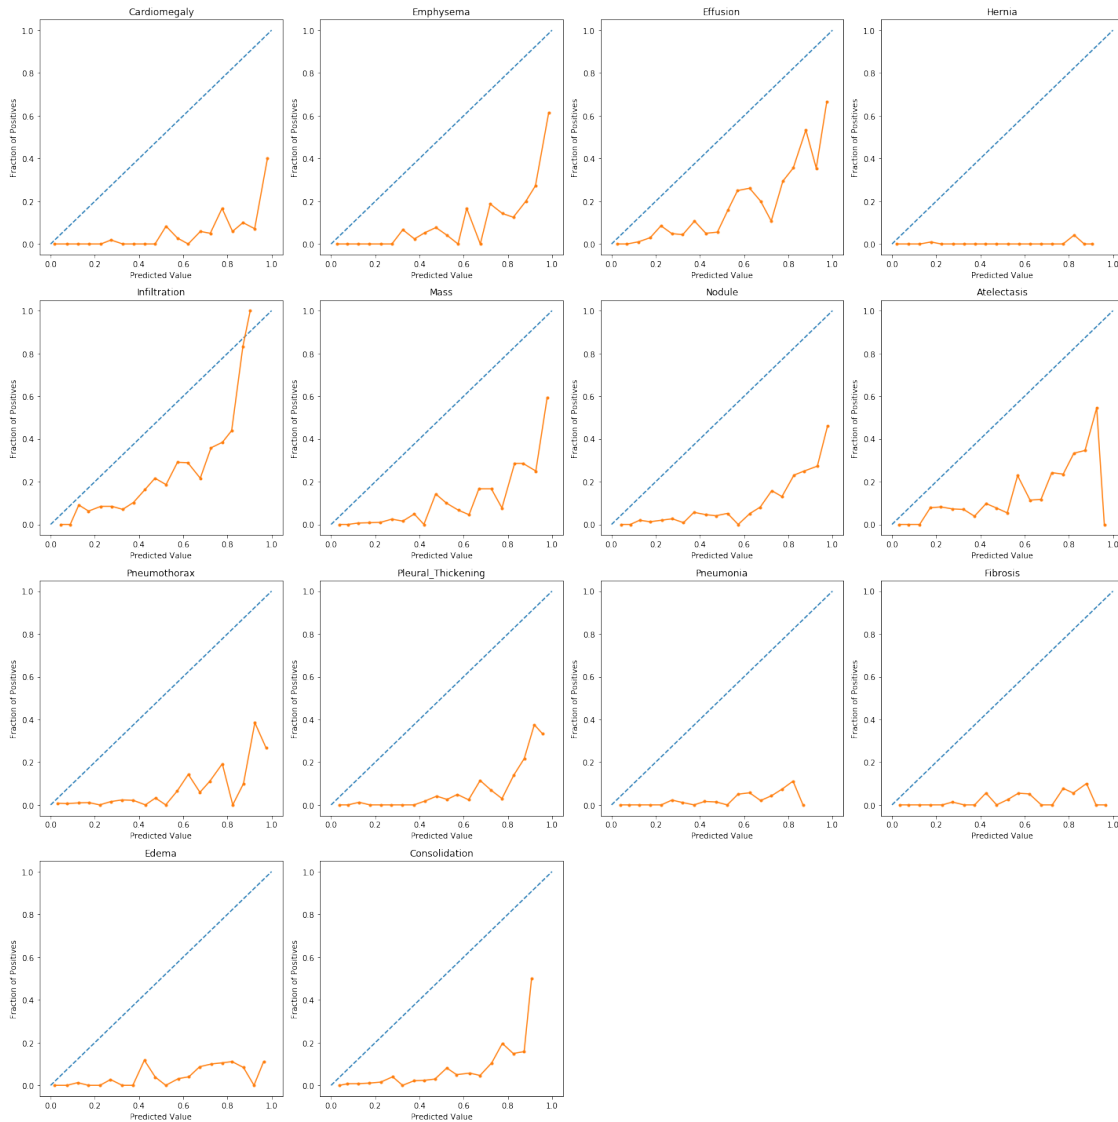
	TP	TN	FP	FN	Accuracy	Prevalence	Sensitivity	\
Cardiomegaly	16	814	169	1	0.83	0.017	0.941	
Emphysema	20	869	103	8	0.889	0.028	0.714	
Effusion	99	690	196	15	0.789	0.114	0.868	
Hernia	1	743	255	1	0.744	0.002	0.5	
Infiltration	114	543	265	78	0.657	0.192	0.594	
Mass	40	789	158	13	0.829	0.053	0.755	
Nodule	28	731	220	21	0.759	0.049	0.571	
Atelectasis	64	657	249	30	0.721	0.094	0.681	

Pneumothorax	24	785	183	8	0.809	0.032	0.75
Pleural_Thickening	24	713	259	4	0.737	0.028	0.857
Pneumonia	14	661	320	5	0.675	0.019	0.737
Fibrosis	10	725	261	4	0.735	0.014	0.714
Edema	15	767	213	5	0.782	0.02	0.75
Consolidation	36	658	297	9	0.694	0.045	0.8

	Specificity	PPV	NPV	AUC	F1	Threshold
Cardiomegaly	0.828	0.086	0.999	0.933	0.158	0.5
Emphysema	0.894	0.163	0.991	0.935	0.265	0.5
Effusion	0.779	0.336	0.979	0.891	0.484	0.5
Hernia	0.744	0.004	0.999	0.644	0.008	0.5
Infiltration	0.672	0.301	0.874	0.696	0.399	0.5
Mass	0.833	0.202	0.984	0.888	0.319	0.5
Nodule	0.769	0.113	0.972	0.745	0.189	0.5
Atelectasis	0.725	0.204	0.956	0.781	0.314	0.5
Pneumothorax	0.811	0.116	0.99	0.826	0.201	0.5
Pleural_Thickening	0.734	0.085	0.994	0.868	0.154	0.5
Pneumonia	0.674	0.042	0.992	0.762	0.079	0.5
Fibrosis	0.735	0.037	0.995	0.801	0.07	0.5
Edema	0.783	0.066	0.994	0.856	0.121	0.5
Consolidation	0.689	0.108	0.987	0.799	0.19	0.5

```
In [30]: from sklearn.calibration import calibration_curve
def plot_calibration_curve(y, pred):
    plt.figure(figsize=(20, 20))
    for i in range(len(class_labels)):
        plt.subplot(4, 4, i + 1)
        fraction_of_positives, mean_predicted_value = calibration_curve(y[:,i], pred)
        plt.plot([0, 1], [0, 1], linestyle='--')
        plt.plot(mean_predicted_value, fraction_of_positives, marker='.')
        plt.xlabel("Predicted Value")
        plt.ylabel("Fraction of Positives")
        plt.title(class_labels[i])
    plt.tight_layout()
    plt.show()
```

```
In [31]: plot_calibration_curve(y, pred)
```

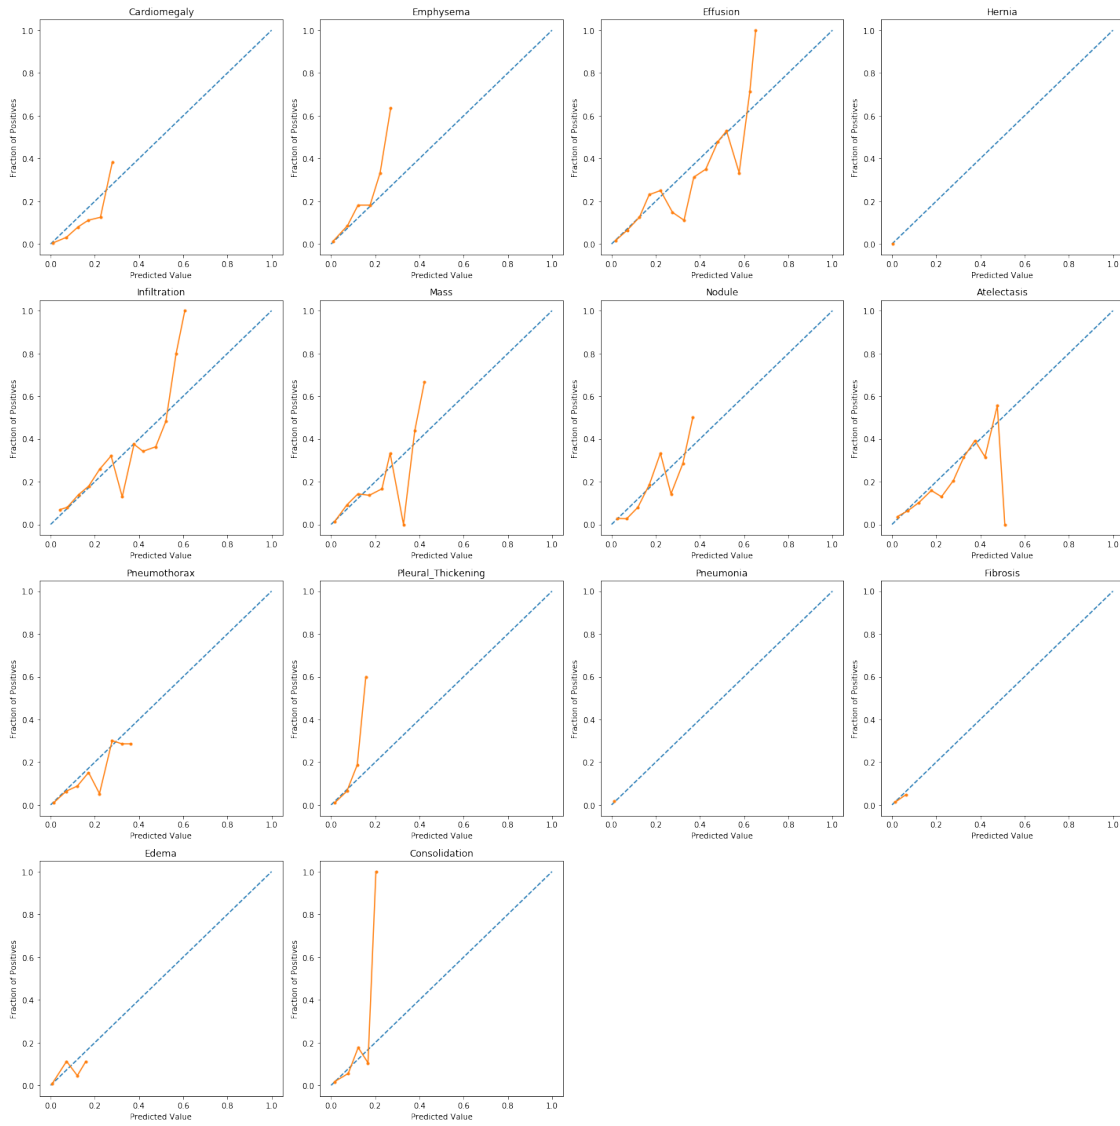


```
In [32]: from sklearn.linear_model import LogisticRegression as LR
```

```
y_train = train_results[class_labels].values
pred_train = train_results[pred_labels].values
pred_calibrated = np.zeros_like(pred)

for i in range(len(class_labels)):
    lr = LR(solver='liblinear', max_iter=10000)
    lr.fit(pred_train[:, i].reshape(-1, 1), y_train[:, i])
    pred_calibrated[:, i] = lr.predict_proba(pred[:, i].reshape(-1, 1))[:, 1]
```

```
In [33]: plot_calibration_curve(y[:, :], pred_calibrated)
```



*****END*****

In []: Project By: Aaditya Balakrishnan